



KUBERNATIC

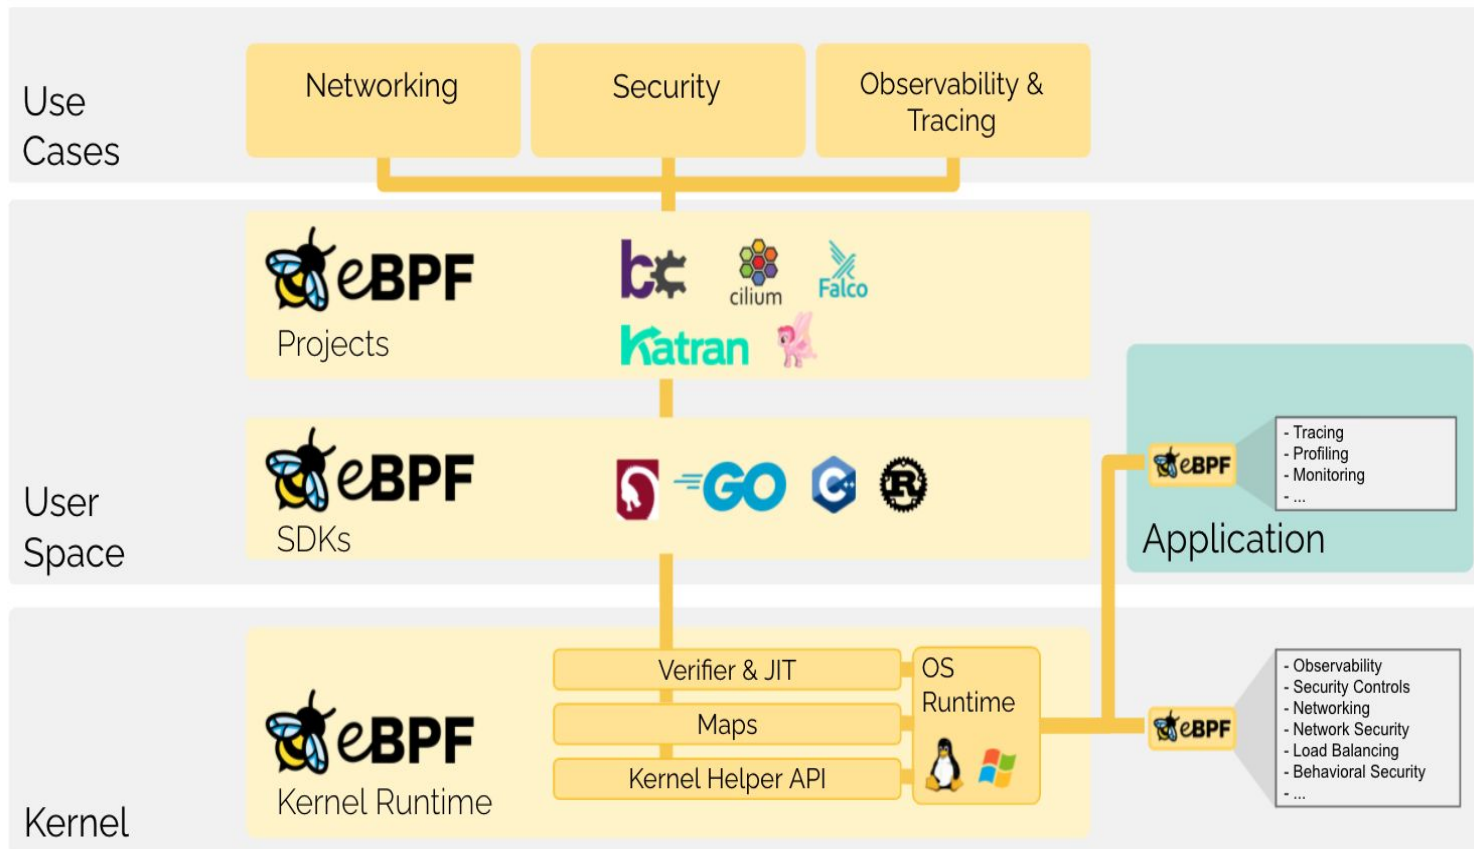
Hackathon - Berlin 2022

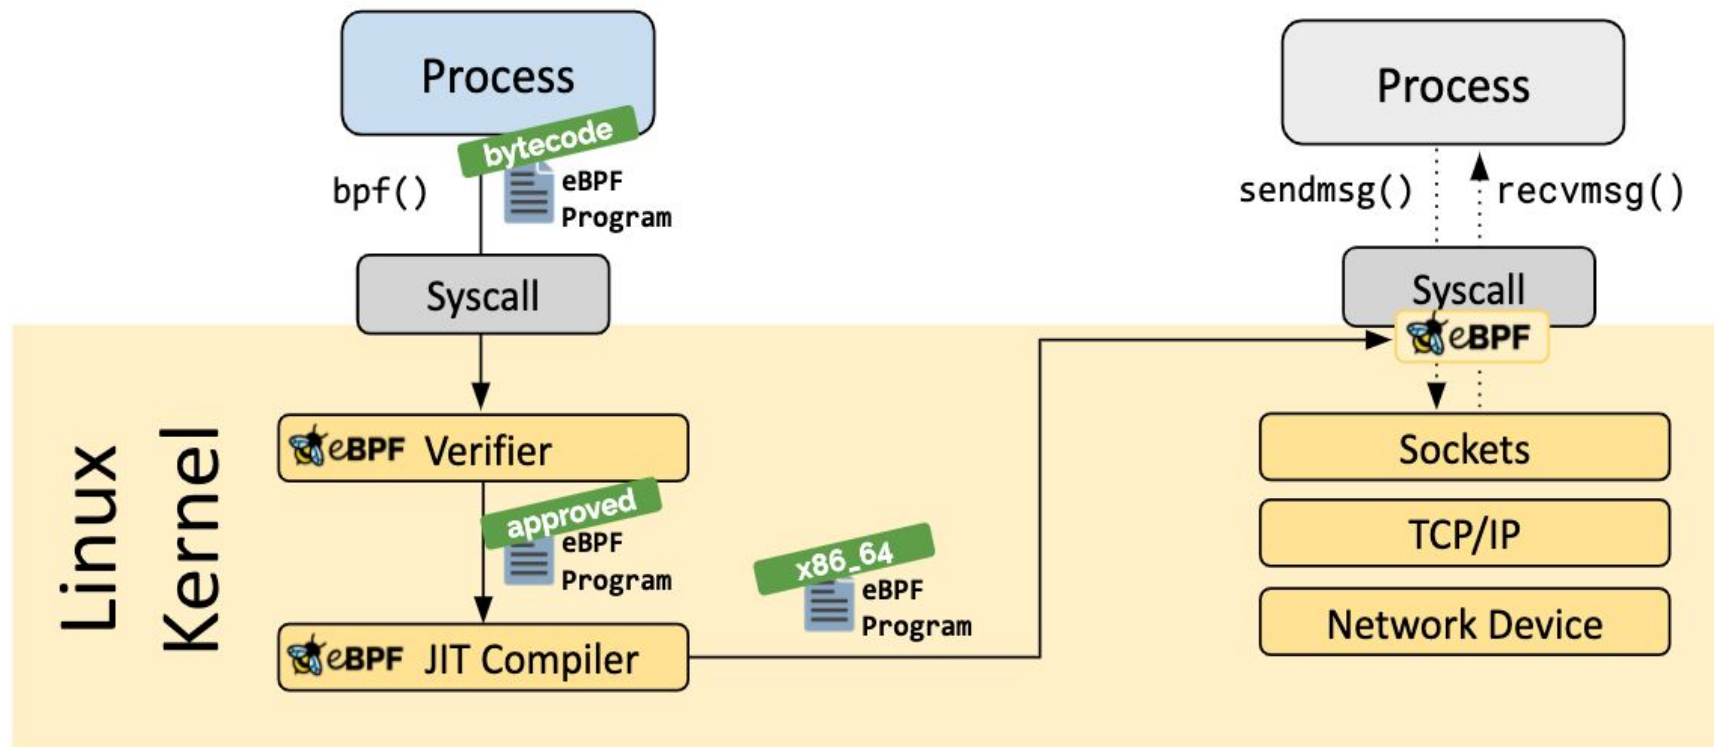
**Learn some eBPF for
greater good**



eBPF

theoretical introduction





uprobe

writing runtime code execution agent from scratch

```
import (  
    "fmt"  
    "time"  
)  
  
func easyToFindFunctionName(arg uint32) {  
    fmt.Println(arg)  
}  
  
func main() {  
    t1, t2 := time.NewTicker(time.Second * 3), time.NewTicker(time.Second * 5)  
    for {  
        select {  
        case <-t1.C:  
            easyToFindFunctionName(1)  
        case <-t2.C:  
            easyToFindFunctionName(2)  
        }  
    }  
}
```

```
import (  
    "fmt"  
    "time"  
)  
  
func easyToFindFunctionName(arg uint32) {  
    fmt.Println(arg)  
}  
  
func main() {  
    t1, t2 := time.NewTicker(time.Second * 3), time.NewTicker(time.Second * 5)  
    for {  
        select {  
        case <-t1.C:  
            easyToFindFunctionName(1)  
        case <-t2.C:  
            easyToFindFunctionName(2)  
        }  
    }  
}
```

```
import (  
    "fmt"  
    "time"  
)  
  
func easyToFindFunctionName(arg uint32) {  
    fmt.Println(arg)  
}  
  
func main() {  
    t1, t2 := time.NewTicker(time.Second * 3), time.NewTicker(time.Second * 5)  
    for {  
        select {  
        case <-t1.C:  
            easyToFindFunctionName(1)  
        case <-t2.C:  
            easyToFindFunctionName(2)  
        }  
    }  
}
```



```
import (  
    "fmt"  
    "time"  
)
```

```
func easyToFindFunctionName(arg uint32) {  
    fmt.Println(arg)  
}
```

```
fun $ go build -gcflags '-N -l' -o testbin ./main.go  
    t1, t2 := time.NewTicker(time.Second * 5), time.NewTicker(time.Second * 5)  
    for {  
        select {  
        case <-t1.C:  
            easyToFindFunctionName(1)  
        case <-t2.C:  
            easyToFindFunctionName(2)  
        }  
    }  
}
```



```
import (  
    "github.com/cilium/ebpf"  
    "github.com/cilium/ebpf/link"  
)  
  
func main() {  
    // Load ebpf byte code and extract maps and programs  
    ebpf.LoadAndAssign(obj, opts)  
  
    // open executable to be instrumented  
    ex, err := link.OpenExecutable("/bin/testbin")  
  
    // register uprobe on a particular symbol in a binary  
    up, err := ex.Uprobe("easyToFindFunctionName", objs.UprobeTestbinTest)  
  
    // listen on events from kernel through perf event map  
    // BPF_MAP_TYPE_PERF_EVENT_ARRAY  
    for{  
        fmt.Println(ts, event.Pid, symbol, event.Arg)  
    }  
}
```

```
import (  
    "github.com/cilium/ebpf"  
    "github.com/cilium/ebpf/link"  
)  
  
func main() {  
    // Load ebpf byte code and extract maps and programs  
    ebpf.LoadAndAssign(obj, opts)  
  
    // open executable to be instrumented  
    ex, err := link.OpenExecutable("/bin/testbin")  
  
    // register uprobe on a particular symbol in a binary  
    up, err := ex.Uprobe("easyToFindFunctionName", objs.UprobeTestbinTest)  
  
    // listen on events from kernel through perf event map  
    // BPF_MAP_TYPE_PERF_EVENT_ARRAY  
    for{  
        fmt.Println(ts, event.Pid, symbol, event.Arg)  
    }  
}
```



```
import (  
    "github.com/cilium/ebpf"  
    "github.com/cilium/ebpf/link"  
)  
  
func main() {  
    // Load ebpf byte code and extract maps and programs  
    ebpf.LoadAndAssign(obj, opts)  
  
    // open executable to be instrumented  
    ex, err := link.OpenExecutable("/bin/testbin")  
  
    // register uprobe on a particular symbol in a binary  
    up, err := ex.Uprobe("easyToFindFunctionName", objs.UprobeTestbinTest)  
  
    // listen on events from kernel through perf event map  
    // BPF_MAP_TYPE_PERF_EVENT_ARRAY  
    for{  
        fmt.Println(ts, event.Pid, symbol, event.Arg)  
    }  
}
```



```
import (  
    "github.com/cilium/ebpf"  
    "github.com/cilium/ebpf/link"  
)  
  
func main() {  
    // Load ebpf byte code and extract maps and programs  
    ebpf.LoadAndAssign(obj, opts)  
  
    // open executable to be instrumented  
    ex, err := link.OpenExecutable("/bin/testbin")  
  
    // register uprobe on a particular symbol in a binary  
    up, err := ex.Uprobe("easyToFindFunctionName", objs.UprobeTestbinTest)  
  
    // listen on events from kernel through perf event map  
    // BPF_MAP_TYPE_PERF_EVENT_ARRAY  
    for{  
        fmt.Println(ts, event.Pid, symbol, event.Arg)  
    }  
}
```

```
import (  
    "github.com/cilium/ebpf"  
    "github.com/cilium/ebpf/link"  
)  
  
func main() {  
    // Load ebpf byte code and extract maps and programs  
    ebpf.LoadAndAssign(obj, opts)  
  
    // open executable to be instrumented  
    ex, err := link.OpenExecutable("/bin/testbin")  
  
    // register uprobe on a particular symbol in a binary  
    up, err := ex.Uprobe("easyToFindFunctionName", objs.UprobeTestbinTest)  
  
    // listen on events from kernel through perf event map  
    // BPF_MAP_TYPE_PERF_EVENT_ARRAY  
    for{  
        fmt.Println(ts, event.Pid, symbol, event.Arg)  
    }  
}
```

```
#include "bpf_tracing.h"

struct event {
    u32 pid;
    u32 arg;
};

struct {
    __uint(type, BPF_MAP_TYPE_PERF_EVENT_ARRAY);
} events SEC(".maps");

SEC("uprobe/testbin_test")
int uprobe_testbin_test(struct pt_regs *ctx) {
    struct event event;
    bpf_probe_read(&event.arg, sizeof(event.arg), (void*)PT_REGS_SP(ctx)+8);
    if (event.arg == 2) {
        event.pid = bpf_get_current_pid_tgid();
        bpf_perf_event_output(ctx, &events, BPF_F_CURRENT_CPU, &event, sizeof(event));
    }
    return 0;
}
```



```
#include "bpf_tracing.h"
```

```
struct event {  
    u32 pid;  
    u32 arg;  
};
```

```
struct {  
    __uint(type, BPF_MAP_TYPE_PERF_EVENT_ARRAY);  
} events SEC(".maps");
```

```
SEC("uprobe/testbin_test")
```

```
int uprobe_testbin_test(struct pt_regs *ctx) {  
    struct event event;  
    bpf_probe_read(&event.arg, sizeof(event.arg), (void*)PT_REGS_SP(ctx)+8);  
    if (event.arg == 2) {  
        event.pid = bpf_get_current_pid_tgid();  
        bpf_perf_event_output(ctx, &events, BPF_F_CURRENT_CPU, &event, sizeof(event));  
    }  
    return 0;  
}
```



```
#include "bpf_tracing.h"

struct event {
    u32 pid;
    u32 arg;
};

struct {
    __uint(type, BPF_MAP_TYPE_PERF_EVENT_ARRAY);
} events SEC(".maps");

SEC("uprobe/testbin_test")
int uprobe_testbin_test(struct pt_regs *ctx) {
    struct event event;
    bpf_probe_read(&event.arg, sizeof(event.arg), (void*)PT_REGS_SP(ctx)+8);
    if (event.arg == 2) {
        event.pid = bpf_get_current_pid_tgid();
        bpf_perf_event_output(ctx, &events, BPF_F_CURRENT_CPU, &event, sizeof(event));
    }
    return 0;
}
```

```
#include "bpf_tracing.h"

struct event {
    u32 pid;
    u32 arg;
};

struct {
    __uint(type, BPF_MAP_TYPE_PERF_EVENT_ARRAY);
} events SEC(".maps");

SEC("uprobe/testbin_test")
int uprobe_testbin_test(struct pt_regs *ctx) {
    struct event event;
    bpf_probe_read(&event.arg, sizeof(event.arg), (void*)PT_REGS_SP(ctx)+8);
    if (event.arg == 2) {
        event.pid = bpf_get_current_pid_tgid();
        bpf_perf_event_output(ctx, &events, BPF_F_CURRENT_CPU, &event, sizeof(event));
    }
    return 0;
}
```

```
#include "bpf_tracing.h"

struct event {
    u32 pid;
    u32 arg;
};

struct {
    __uint(type, BPF_MAP_TYPE_PERF_EVENT_ARRAY);
} events SEC(".maps");

SEC("uprobe/testbin_test")
int uprobe_testbin_test(struct pt_regs *ctx) {
    struct event event;
    bpf_probe_read(&event.arg, sizeof(event.arg), (void*)PT_REGS_SP(ctx)+8);
    if (event.arg == 2) {
        event.pid = bpf_get_current_pid_tgid();
        bpf_perf_event_output(ctx, &events, BPF_F_CURRENT_CPU, &event, sizeof(event));
    }
    return 0;
}
```

```
#include "bpf_tracing.h"

struct event {
    u32 pid;
    u32 arg;
};

struct {
    __uint(type, BPF_MAP_TYPE_PERF_EVENT_ARRAY);
} events SEC(".maps");

SEC("uprobe/testbin_test")
int uprobe_testbin_test(struct pt_regs *ctx) {
    struct event event;
    bpf_probe_read(&event.arg, sizeof(event.arg), (void*)PT_REGS_SP(ctx)+8);
    if (event.arg == 2) {
        event.pid = bpf_get_current_pid_tgid();
        bpf_perf_event_output(ctx, &events, BPF_F_CURRENT_CPU, &event, sizeof(event));
    }
    return 0;
}
```

```
#include "bpf_tracing.h"

struct event {
    u32 pid;
    u32 arg;
};

struct {
    __uint(type, BPF_MAP_TYPE_PERF_EVENT_ARRAY)
} events SEC(".maps");

SEC("uprobe/testbin/test")
int uprobe_testbin_test(struct pt_regs *ctx) {
    struct event event;
    bpf_probe_read(&event.arg, sizeof(event.arg), (void*)PT_REGS_SP(ctx)+8);
    if (event.arg == 2) {
        event.pid = bpf_get_current_pid_tgid();
        bpf_perf_event_output(ctx, &events, BPF_F_CURRENT_CPU, &event, sizeof(event));
    }
    return 0;
}
```

Disclaimer

```
#include <linux/bpf.h>
...
struct event {
    u32 pid;
    u32 tgid;
};

SEC("uprobe")
int uprobe(struct event *event) {
    bpf_probe_read(&event->arg, sizeof(event->arg), (void*)PT_REGS_SP(ctx)+8);
    if (event->arg == 2) {
        event->pid = bpf_get_current_pid_tgid();
        bpf_perf_event_output(ctx, &events, BPF_F_CURRENT_CPU, &event, sizeof(event));
    }
    return 0;
}
```

```
TEXT main.easyToFindFunctionName(SB) /ebpf-hackathon/uprobe_call_detect/test_bin.go
test_bin.go:11 0x48e700 493b6610 CMPQ 0x10(R14), SP
test_bin.go:11 0x48e70a 4883ec58 SUBQ $0x58, SP
test_bin.go:11 0x48e70e 48896c2450 MOVQ BP, 0x50(SP)
test_bin.go:12 0x48e76c 488b442420 MOVQ 0x20(SP), AX
...
test_bin.go:12 0x48e794 e867aaffff CALL fmt.Println(SB)
...
test_bin.go:15 0x48e7f4 ebca JMP main.easyToFindFunctionName(SB)

TEXT main.main(SB) /ebpf-hackathon/uprobe_call_detect/test_bin.go
...
test_bin.go:27 0x48e8d2 b802000000 MOVL $0x2, 0x18(SP)
test_bin.go:27 0x48e8d7 e8e4feffff CALL main.easyToFindFunctionName(SB)
test_bin.go:26 0x48e8dc eb0c JMP 0x48e8ea
test_bin.go:25 0x48e8de b801000000 MOVL $0x1, 0x18(SP)
test_bin.go:25 0x48e8e3 e8d8feffff CALL main.easyToFindFunctionName(SB)
...
```

```
#include <linux/bpf.h>
struct event {
    u32 pid;
    u32 tgid;
};

struct event event;

SEC("uprobe")
int uprobe(struct pt_regs *regs) {
    struct event event;
    bpf_probe_read(&event.arg, sizeof(event.arg), (void*)PT_REGS_SP(ctx)+8);
    if (event.arg == 2) {
        event.pid = bpf_get_current_pid_tgid();
        bpf_perf_event_output(ctx, &events, BPF_F_CURRENT_CPU, &event, sizeof(event));
    }
    return 0;
}
```

```
TEXT main.easyToFindFunctionName(SB) /ebpf-hackathon/uprobe_call_detect/test_bin.go
test_bin.go:11 0x48e700 493b6610 CMPQ 0x10(R14), SP
test_bin.go:11 0x48e70a 4883ec58 SUBQ $0x58, SP
test_bin.go:11 0x48e70e 48896c2450 MOVQ BP, 0x50(SP)
test_bin.go:12 0x48e76c 488b442420 MOVQ 0x20(SP), AX
...
test_bin.go:12 0x48e794 e867aaffff CALL fmt.Println(SB)
...
test_bin.go:15 0x48e7f4 ebca JMP main.easyToFindFunctionName(SB)

TEXT main.main(SB) /ebpf-hackathon/uprobe_call_detect/test_bin.go
...
test_bin.go:27 0x48e8d2 b802000000 MOVL $0x2, 0x18(SP)
test_bin.go:27 0x48e8d7 e8e4feffff CALL main.easyToFindFunctionName(SB)
test_bin.go:26 0x48e8dc eb0c JMP 0x48e8ea
test_bin.go:25 0x48e8de b801000000 MOVL $0x1, 0x18(SP)
test_bin.go:25 0x48e8e3 e8d8feffff CALL main.easyToFindFunctionName(SB)
...

```

```
#include <linux/bpf.h>
...
struct event {
    u32 pid;
    u32 tgid;
};

SEC("uprobe")
int uprobe_handler(struct event *event) {
    TEXT main.easyToFindFunctionName(SB) /ebpf-hackathon/uprobe_call_detect/test_bin.go
    test_bin.go:11 0x48e700 493b6610 CMPQ 0x10(R14), SP
    test_bin.go:11 0x48e70a 4883ec58 SUBQ $0x58, SP
    test_bin.go:11 0x48e70e 48896c2450 MOVQ BP, 0x50(SP)
    test_bin.go:12 0x48e76c 488b442420 MOVQ 0x20(SP), AX
    ...
    test_bin.go:12 0x48e794 e867aaffff CALL fmt.Println(SB)
    ...
    test_bin.go:15 0x48e7f4 ebca JMP main.easyToFindFunctionName(SB)

    TEXT main.main(SB) /ebpf-hackathon/uprobe_call_detect/test_bin.go
    ...
    test_bin.go:27 0x48e8d2 b802000000 MOVL $0x2, 0x18(SP)
    test_bin.go:27 0x48e8d7 e8e4feffff CALL main.easyToFindFunctionName(SB)
    test_bin.go:26 0x48e8dc eb0c JMP 0x48e8ea
    test_bin.go:25 0x48e8de b801000000 MOVL $0x1, 0x18(SP)
    test_bin.go:25 0x48e8e3 e8d8feffff CALL main.easyToFindFunctionName(SB)
    ...
}
```

```
    struct event event;
```

```
    bpf_probe_read(&event.arg, sizeof(event.arg), (void*)PT_REGS_SP(ctx)+8);
```

```
    if (event.arg == 2) {
```

```
        event.pid = bpf_get_current_pid_tgid();
```

```
        bpf_perf_event_output(ctx, &events, BPF_F_CURRENT_CPU, &event, sizeof(event));
```

```
    }
```

```
    return 0;
```

```
}
```



```
#include <linux/bpf.h>
struct event {
    u32 pid;
    u32 tgid;
};

struct event event;

SEC("uprobe")
int uprobe_handler(struct pt_regs *ctx) {
    struct event event;
    bpf_probe_read(&event.arg, sizeof(event.arg), (void*)PT_REGS_SP(ctx)+8);
    if (event.arg == 2) {
        event.pid = bpf_get_current_pid_tgid();
        bpf_perf_event_output(ctx, &events, BPF_F_CURRENT_CPU, &event, sizeof(event));
    }
    return 0;
}
```

```
TEXT main.easyToFindFunctionName(SB) /ebpf-hackathon/uprobe_call_detect/test_bin.go
test_bin.go:11 0x48e700 493b6610 CMPQ 0x10(R14), SP
test_bin.go:11 0x48e70a 4883ec58 SUBQ $0x58, SP
test_bin.go:11 0x48e70e 48896c2450 MOVQ BP, 0x50(SP)
test_bin.go:12 0x48e76c 488b442420 MOVQ 0x20(SP), AX
...
test_bin.go:12 0x48e794 e867aaffff CALL fmt.Println(SB)
...
test_bin.go:15 0x48e7f4 ebca JMP main.easyToFindFunctionName(SB)

TEXT main.main(SB) /ebpf-hackathon/uprobe_call_detect/test_bin.go
...
test_bin.go:27 0x48e8d2 b802000000 MOVL $0x2, 0x18(SP)
test_bin.go:27 0x48e8d7 e8e4feffff CALL main.easyToFindFunctionName(SB)
test_bin.go:26 0x48e8dc eb0c JMP 0x48e8ea
test_bin.go:25 0x48e8de b801000000 MOVL $0x1, 0x18(SP)
test_bin.go:25 0x48e8e3 e8d8feffff CALL main.easyToFindFunctionName(SB)
...

```

```
#include <stdio.h>
TEXT main.easyToFindFunctionName(SB) /ebpf-hackathon/uprobe_call_detect/test_bin.go
    test_bin.go:11  0x48e700  493b6610  CMPQ 0x10(R14), SP
    test_bin.go:11  0x48e70a  4883ec58  SUBQ $0x58, SP
    test_bin.go:11  0x48e70e  48896c2450 MOVQ BP, 0x50(SP)
    test_bin.go:12  0x48e76c  488b442420 MOVQ 0x20(SP), AX
    ...
    test_bin.go:12  0x48e794  e867aaffff  CALL fmt.Println(SB)
    ...
    test_bin.go:15  0x48e7f4  ebca      JMP main.easyToFindFunctionName(SB)

TEXT main.main(SB) /ebpf-hackathon/uprobe_call_detect/test_bin.go
    test_bin.go:25  0x48e8de  b801000000 MOVL $0x1, 0x18(SP)
    test_bin.go:25  0x48e8e3  e8d8feffff  CALL main.easyToFindFunctionName(SB)
    ...
    struct event event;
    bpf_probe_read(&event.arg, sizeof(event.arg), (void*)PT_REGS_SP(ctx)+8);
    if (event.arg == 2) {
        event.pid = bpf_get_current_pid_tgid();
        bpf_perf_event_output(ctx, &events, BPF_F_CURRENT_CPU, &event, sizeof(event));
    }
    return 0;
}
```

```
$ go build -gcflags '-N -l' -o testbin ./main.go
```

```
SEC("uprobe")
int uprobe(struct pt_regs *ctx) {
    struct event event;
    bpf_probe_read(&event.arg, sizeof(event.arg), (void*)PT_REGS_SP(ctx)+8);
    if (event.arg == 2) {
        event.pid = bpf_get_current_pid_tgid();
        bpf_perf_event_output(ctx, &events, BPF_F_CURRENT_CPU, &event, sizeof(event));
    }
    return 0;
}
```

```
$ k get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
testbin	2/2	2	2	41h

```
$ k get daemonsets
```

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
tracer	9	9	0	9	0	<none>	41h

```
$ k get pods
```

NAME	READY	STATUS	NODE
testbin-764b7bbdd9-h8kcx	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
testbin-764b7bbdd9-nt5hp	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
tracer-7jctq	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
...			

```
$ k get pods
```

NAME	READY	STATUS	NODE
testbin-764b7bbdd9-h8kcx	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
testbin-764b7bbdd9-nt5hp	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
tracer-7jctq	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
...			

```
$ k logs tracer-7jctq
```

```
2022/10/14 07:44:58 instrumenting
/host/io.containerd.snapshotter.v1.overlayfs/snapshots/265/fs/bin/testbin
2022/10/14 07:44:59 Listening for events..
2022/10/14 07:45:01 10536 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:02 10594 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:04 10536 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:05 10594 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:09 10594 main.easyToFindFunctionName argument: 2
```

```
$ k get pods
```

NAME	READY	STATUS	NODE
testbin-764b7bbdd9-h8kcx	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
testbin-764b7bbdd9-nt5hp	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
tracer-7jctq	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
...			

```
$ k logs tracer-7jctq
```

```
2022/10/14 07:44:58 instrumenting
/host/io.containerd.snapshotter.v1.overlayfs/snapshots/265/fs/bin/testbin
2022/10/14 07:44:59 Listening for events..
2022/10/14 07:45:01 10536 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:02 10594 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:04 10536 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:05 10594 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:09 10594 main.easyToFindFunctionName argument: 2
```

```
$ k get pods
```

NAME	READY	STATUS	NODE
testbin-764b7bbdd9-h8kcx	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
testbin-764b7bbdd9-nt5hp	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
tracer-7jctq	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
...			

```
$ k logs tracer-7jctq
```

```
2022/10/14 07:44:58 instrumenting
/host/io.containerd.snapshotter.v1.overlayfs/snapshots/265/fs/bin/testbin
2022/10/14 07:44:59 Listening for events..
2022/10/14 07:45:01 10536 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:02 10594 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:04 10536 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:05 10594 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:09 10594 main.easyToFindFunctionName argument: 2
```

```
$ k get pods
```

NAME	READY	STATUS	NODE
testbin-764b7bbdd9-h8kcx	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
testbin-764b7bbdd9-nt5hp	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
tracer-7jctq	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
...			

```
$ k logs tracer-7jctq
```

```
2022/10/14 07:44:58 instrumenting
/host/io.containerd.snapshotter.v1.overlayfs/snapshots/265/fs/bin/testbin
2022/10/14 07:44:59 Listening for events..
2022/10/14 07:45:01 10536 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:02 10594 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:04 10536 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:05 10594 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:09 10594 main.easyToFindFunctionName argument: 2
```

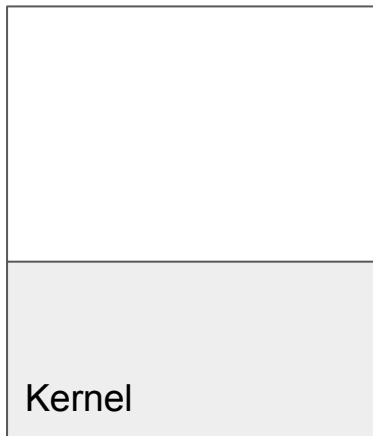
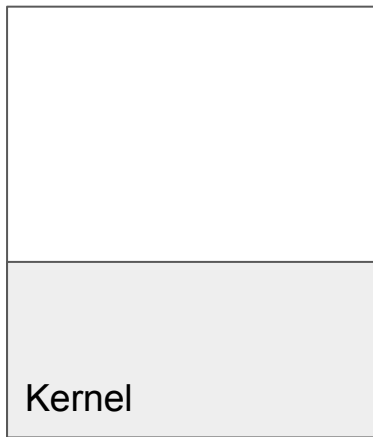
```
$ k get pods
```

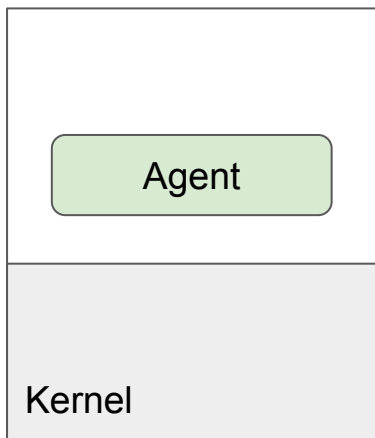
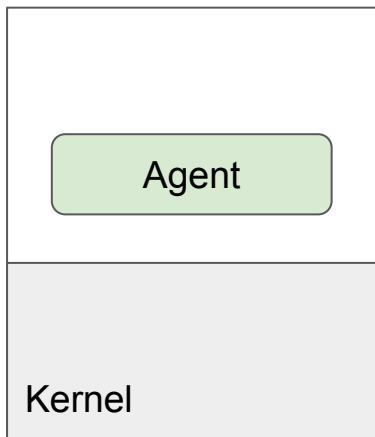
NAME	READY	STATUS	NODE
testbin-764b7bbdd9-h8kcx	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
testbin-764b7bbdd9-nt5hp	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
tracer-7jctq	1/1	Running	gke-kubermatic-dev-cbd686cd-ck6t
...			

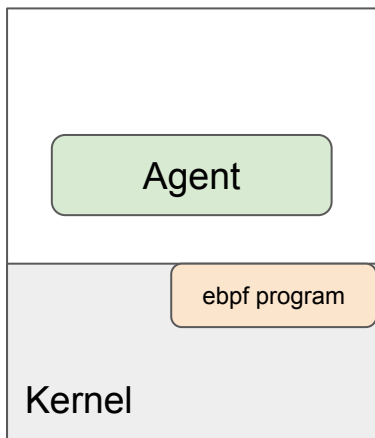
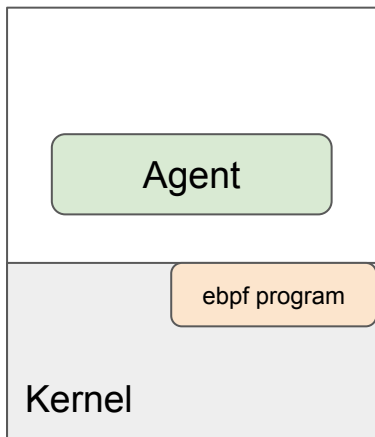
```
$ k logs tracer-7jctq
```

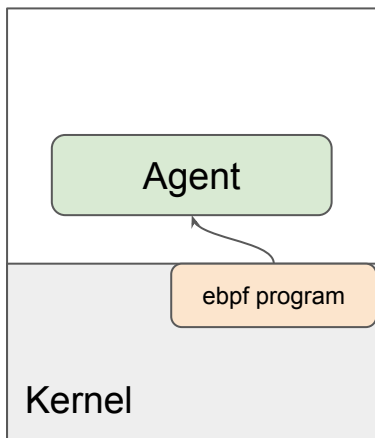
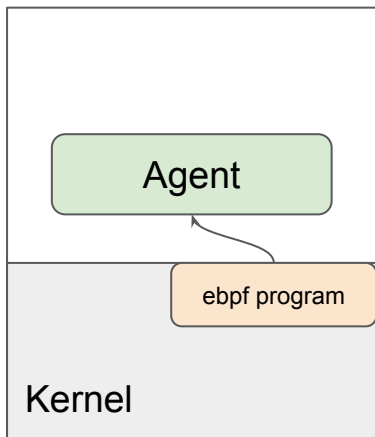
```
2022/10/14 07:44:58 instrumenting
/host/io.containerd.snapshotter.v1.overlayfs/snapshots/265/fs/bin/testbin
2022/10/14 07:44:59 Listening for events..
2022/10/14 07:45:01 10536 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:02 10594 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:04 10536 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:05 10594 main.easyToFindFunctionName argument: 2
2022/10/14 07:45:09 10594 main.easyToFindFunctionName argument: 2
```

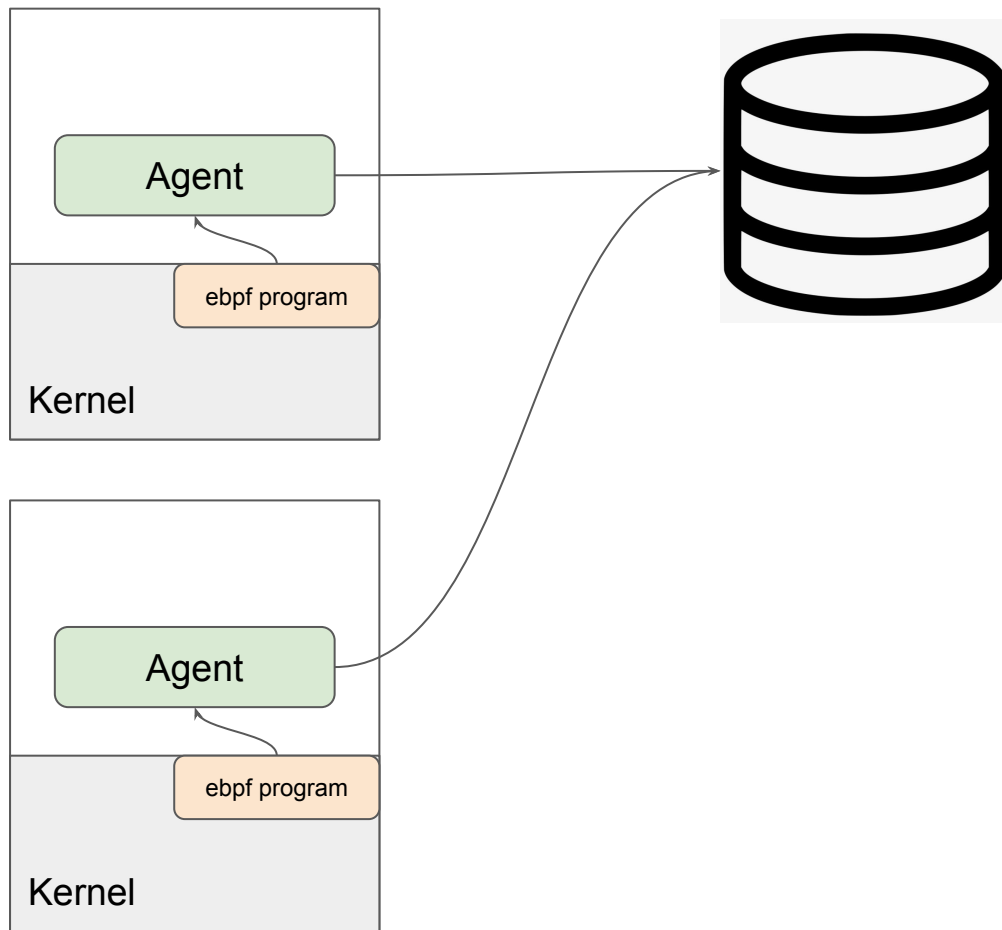

ebpf & Kubernetes







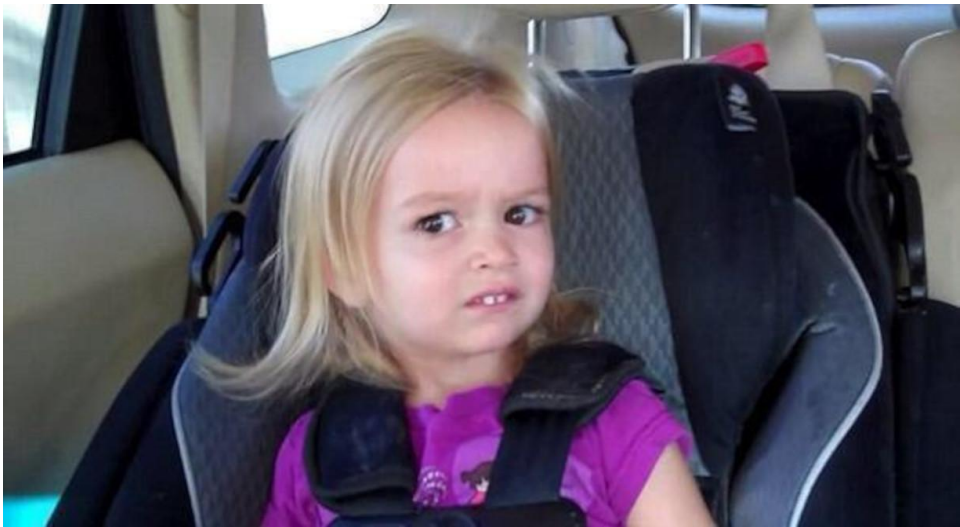






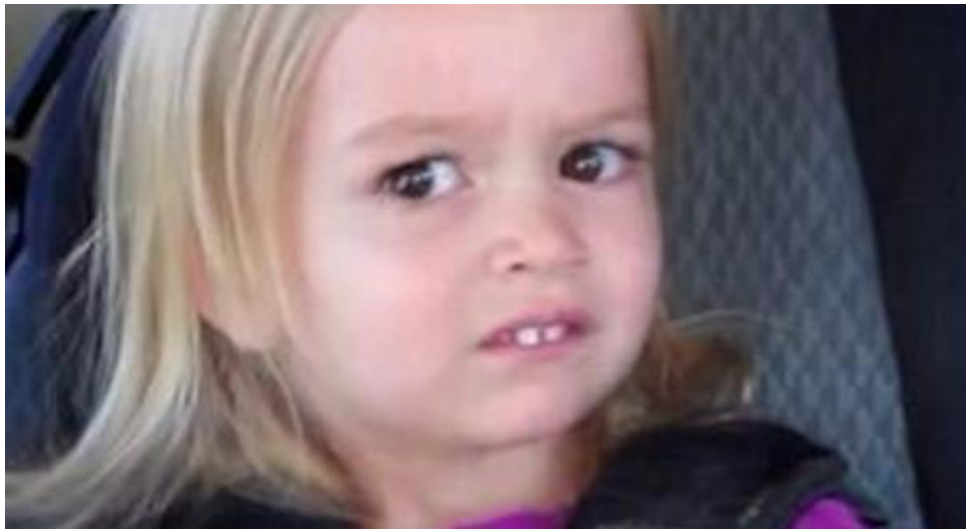
Security?

hostPID: true



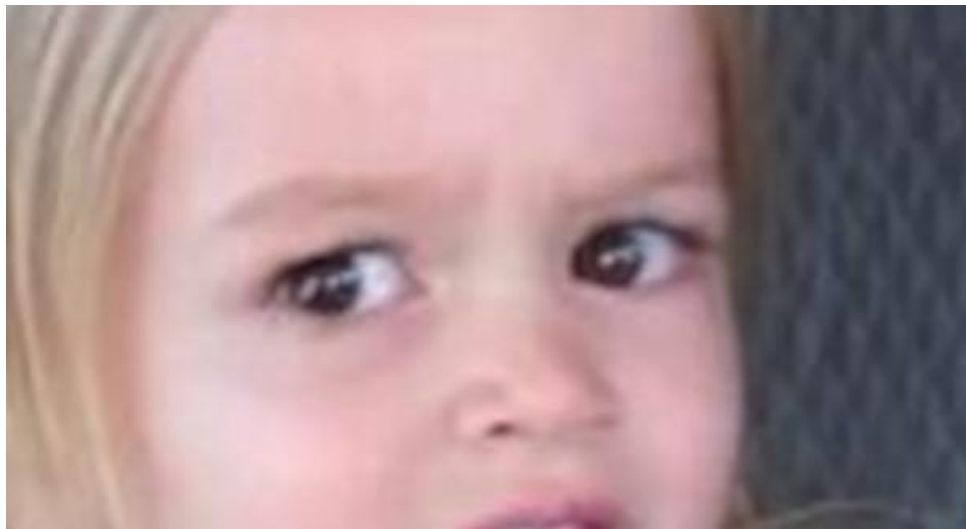
Security?

```
securityContext:  
  privileged: true
```



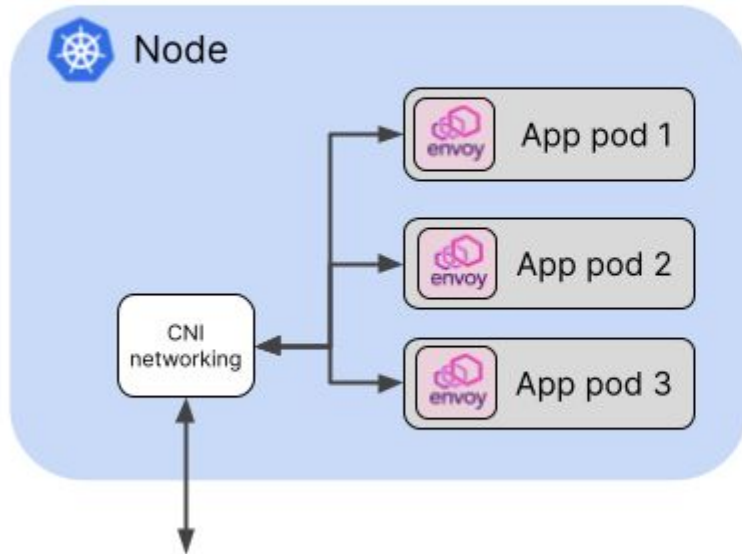
Security?

```
volumes:  
- hostPath:  
  path: /run  
  name: run  
- hostPath:  
  path: /sys/fs/cgroup  
  name: cgroup  
- hostPath:  
  path: /lib/modules  
  name: modules  
- hostPath:  
  path: /sys/fs/bpf  
  name: bpffs  
- hostPath:  
  path: /sys/kernel/debug  
  name: debugfs
```

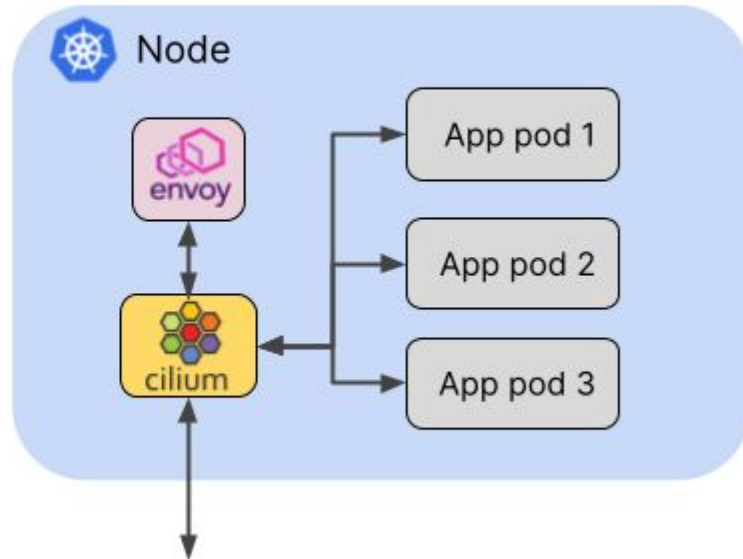


ServiceMesh?

Sidecar Proxy Model



Sidecarless Proxy Model





Parca

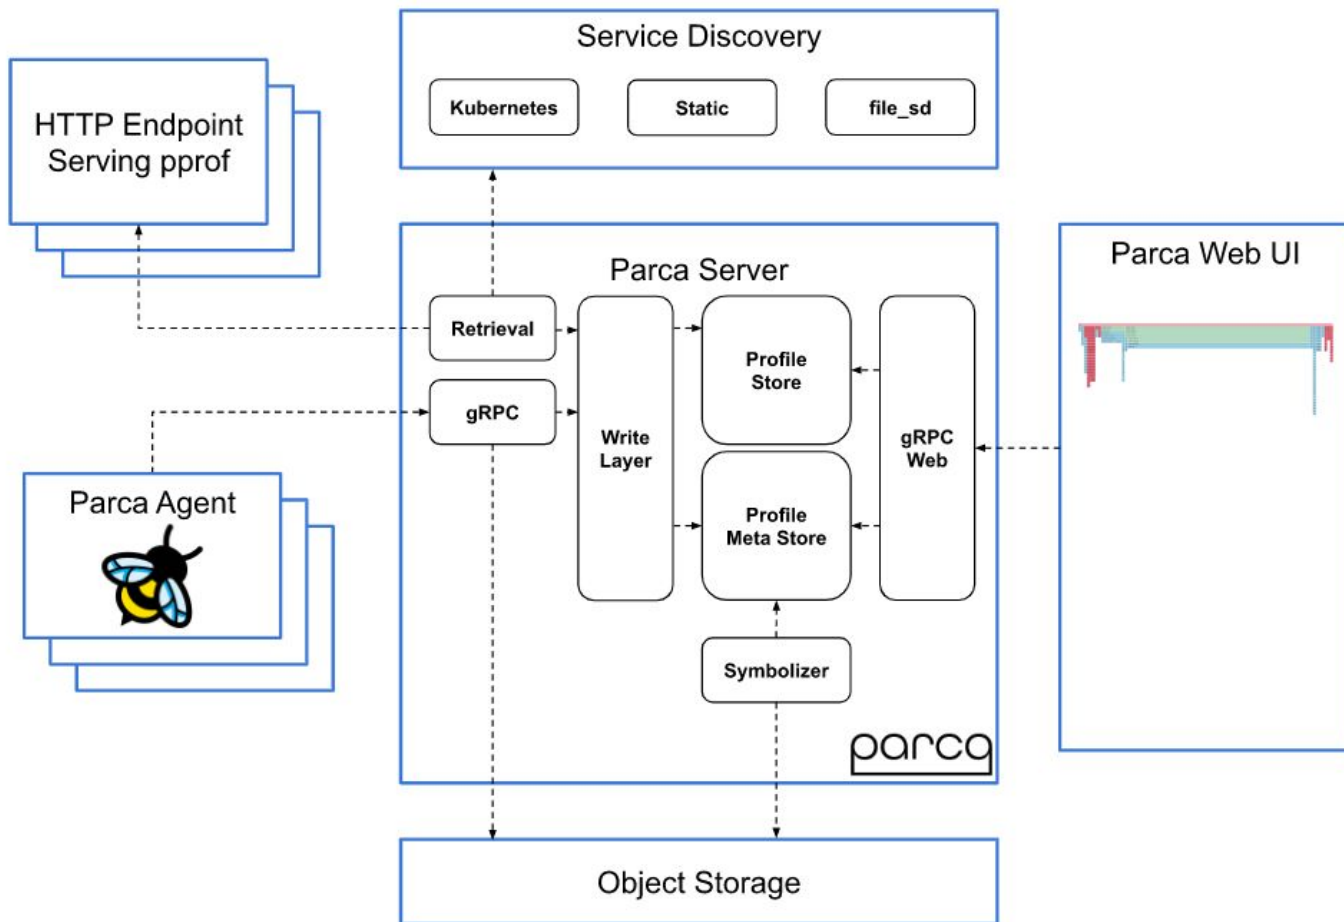


What is Parca?

Continuous profiling for analysis of CPU, memory usage over time, and down to the line number, via eBPF.

Why?

- Save Money
- Improve Performance
- Understand Incidents



root			
[machine-controller]	runtime.goexit		
[machine-controller]	github.com/kubermatic/machine-controller/pkg/controller/machine.NewMachineCollector.func1	[machine-controller]	golang.org/x/ [machine-controller]
[machine-controller]	github.com/kubermatic/machine-controller/pkg/controller/machine.NewMachineCollector.func1.1	[machine-controller]	golang.org/x/ [machine-controller]
[machine-controller]	github.com/kubermatic/machine-controller/pkg/cloudprovider.(*cachingValidationWrapper).SetMetricsForMachines	[machine-controller]	golang.org/x/ [machine-controller]
[machine-controller]	github.com/kubermatic/machine-controller/pkg/cloudprovider/provider/aws.(*provider).SetMetricsForMachines	[machine-controller]	runtime.more [machine-controller]
[machine-controller]	github.com/aws/aws-sdk-go-v2/service/ec2.(*Client).DescribeInstances	[machine-controller]	runtime.newst [machine-controller]
[machine-controller]	github.com/aws/aws-sdk-go-v2/service/ec2.(*Client).invokeOperation	[machine-controller]	runtime.copys [machine-controller]
[machine-controller]	github.com/aws/smithy-go/middleware.(*decoratedHandler).Handle	[machine-controller]	runtime.gentr [machine-controller]
[machine-controller]	github.com/aws/smithy-go/middleware.(*Stack).HandleMiddleware		[machine-controller]
[machine-controller]	github.com/aws/smithy-go/middleware.(*decoratedHandler).Handle		[machine-controller]
[machine-controller]	github.com/aws/smithy-go/middleware.(*InitializeStep).HandleMiddleware		[machine-controller]
[machine-controller]	github.com/aws/smithy-go/middleware.(*decoratedInitializeHandler).HandleInitialize		[machine-controller]
[machine-controller]	github.com/aws/aws-sdk-go-v2/aws/middleware.(*RegisterServiceMetadata).HandleInitialize		[machine-controller]
[machine-controller]	github.com/aws/aws-sdk-go-v2/aws/middleware.RegisterServiceMetadata.HandleInitialize		[machine-controller]
[machine-controller]	github.com/aws/smithy-go/middleware.(*decoratedInitializeHandler).HandleInitialize		[machine-controller]
[machine-controller]	github.com/aws/smithy-go/middleware.(*setLogger).HandleInitialize		[machine-controller]
[machine-controller]	github.com/aws/smithy-go/middleware.(*initializeWrapHandler).HandleInitialize		
[machine-controller]	github.com/aws/smithy-go/middleware.(*decoratedHandler).Handle		
[machine-controller]	github.com/aws/smithy-go/middleware.(*SerializeStep).HandleMiddleware		
[machine-controller]	github.com/aws/smithy-go/middleware.(*decoratedSerializeHandler).HandleSerialize		
[machine-controller]	github.com/aws/aws-sdk-go-v2/service/ec2.(*ResolveEndpoint).HandleSerialize		
[machine-controller]	github.com/aws/smithy-go/middleware.(*decoratedSerializeHandler).HandleSerialize		
[machine-controller]	github.com/aws/aws-sdk-go-v2/service/ec2.(*awsEc2query_serializeOpDescribeInstances).HandleSerialize		
[machine-controller]	github.com/aws/smithy-go/middleware.(*serializeWrapHandler).HandleSerialize		
[machine-controller]	github.com/aws/smithy-go/middleware.(*decoratedHandler).Handle		
[machine-controller]	github.com/aws/smithy-go/middleware.(*BuildStep).HandleMiddleware		
[machine-controller]	github.com/aws/smithy-go/middleware.(*decoratedBuildHandler).HandleBuild		
[machine-controller]	github.com/aws/aws-sdk-go-v2/aws/middleware.(*ClientRequestID).HandleBuild		
[machine-controller]	github.com/aws/aws-sdk-go-v2/aws/middleware.ClientRequestID.HandleBuild		

Our conclusions

Positive Points

- Does not require code instrumentation
- Easy to install
- Overhead is little - we did not crash DEV env
- Nice UI

Negative Points

- UI/server buggy???
- Interpretation of Icicle Graphs is not intuitive
- Currently only CPU metrics supported
- Installation via DaemonSet probably not doable on some projects because of security



Kepler

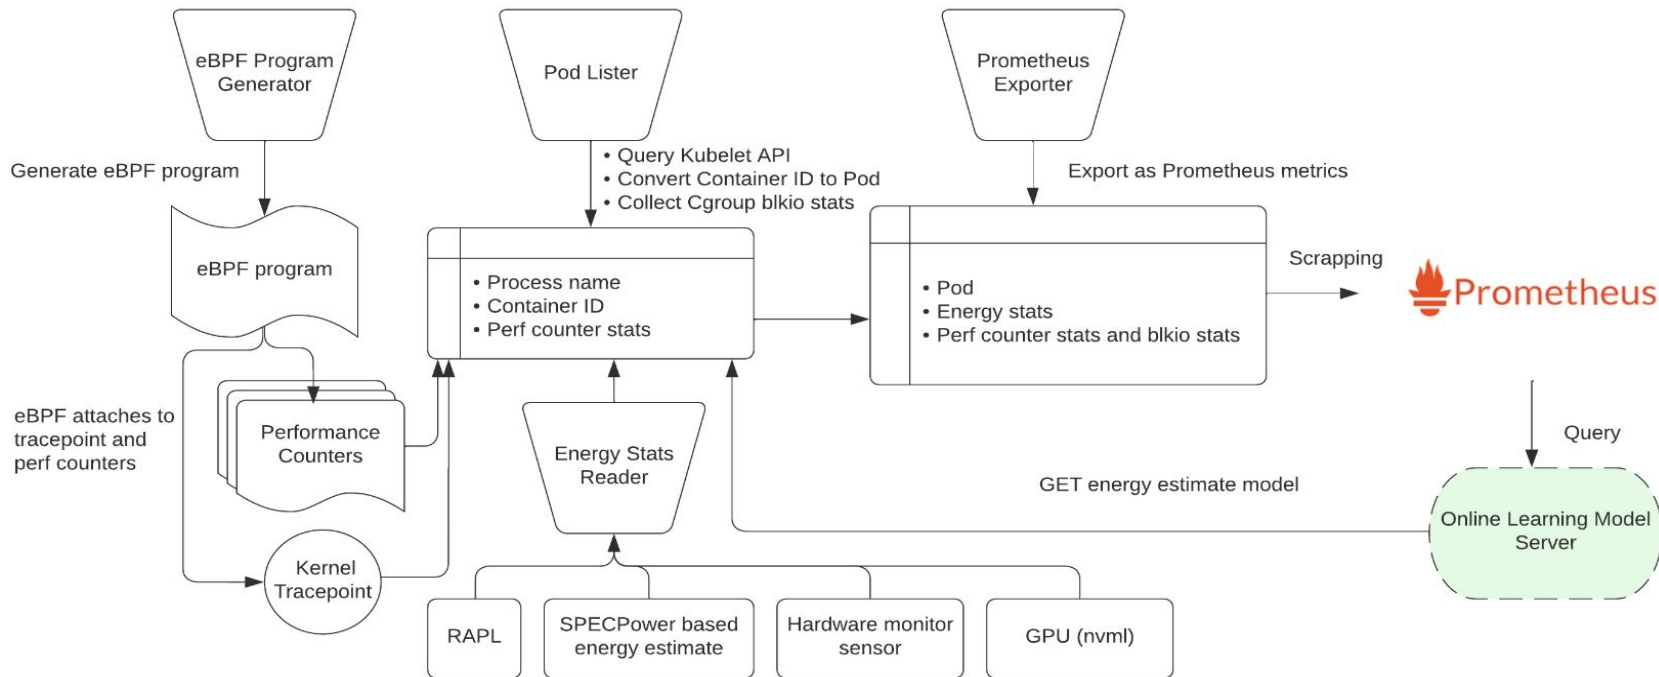
What is Kepler?

Kepler uses eBPF to probe energy related system stats and exports as Prometheus metrics

Why?

- Energy will get more expensive
- Because we can ;)

Kepler: Kubernetes-based Efficient Power Level Exporter



Our conclusions

Positive Points

- Prometheus Metric Scheme looks promising
- Autoscaling based on energy costs?

Negative Points

- It does not work on
 - Ubuntu
 - Rocky Linux
 - CentOS
 - FlatCar
 - RHEL
 - Amazon Linux
- Some kernel module was missing, after installation it still did not work ([github issue](#))